

Annual Review of Economics

The Marginal Propensity to Consume in Heterogeneous Agent Models

Greg Kaplan^{1,3} and Giovanni L. Violante^{2,3,4,5,6,7}

¹Department of Economics, University of Chicago, Chicago, Illinois, USA;
email: gkaplan@uchicago.edu

²Department of Economics, Princeton University, Princeton, New Jersey, USA;
email: violante@princeton.edu

³National Bureau of Economic Research, Cambridge, Massachusetts, USA

⁴Center for Economic Behavior and Inequality, University of Copenhagen, Copenhagen, Denmark

⁵Centre for Economic Policy Research, London, United Kingdom

⁶Institute for Fiscal Studies, London, United Kingdom

⁷Institute of Labor Economics (IZA), Bonn, Germany

Annu. Rev. Econ. 2022. 14:747–75

The *Annual Review of Economics* is online at
economics.annualreviews.org

<https://doi.org/10.1146/annurev-economics-080217-053444>

Copyright © 2022 by Annual Reviews.
All rights reserved

JEL codes: D15, D31, D52, E21, E62, E71, G51

Keywords

borrowing constraints, consumption, hand-to-mouth, heterogeneity, income risk, liquidity, marginal propensity to consume, market incompleteness, precautionary saving, wealth distribution

Abstract

What model features and calibration strategies yield a large average marginal propensity to consume (MPC) in heterogeneous agent models? Through a systematic investigation of models with different preferences, dimensions of ex-ante heterogeneity, income processes, and asset structures, we show that the most important factor is the share and type of hand-to-mouth households. One-asset models either feature a trade-off between a high average MPC and a realistic level of aggregate wealth or generate an excessively polarized wealth distribution that vastly understates the wealth held by households in the middle of the distribution. Two-asset models that include both liquid and illiquid assets can resolve this tension with a large enough gap between liquid and illiquid returns. We discuss how such return differential can be justified from the perspective of theory and data.

ANNUAL REVIEWS CONNECT

www.annualreviews.org

- Download figures
- Navigate cited references
- Keyword search
- Explore related articles
- Share via email or social media

1. INTRODUCTION

The marginal propensity to consume (MPC) is the fraction of a small, unanticipated one-time windfall that a household spends within a given time period. In this article we conduct a systematic investigation of the size and determinants of the aggregate MPC in heterogeneous agent incomplete-markets models—models whose key features are uninsurable idiosyncratic income risk, a precautionary saving motive, and an endogenous wealth distribution.

The MPC is a central concept in modern macroeconomics. In versions of these models with nominal rigidities, it determines the size of fiscal multipliers (Kaplan & Violante 2014, Auclert et al. 2018), the transmission mechanism of monetary policy (Kaplan et al. 2018, Auclert 2019), the amplification of aggregate shocks (Werning 2015, Patterson 2019, Bilbiie 2020, Kekre & Lenel 2021), and portfolio choice between risky and safe assets (Kimball 1990). For these reasons, the usefulness of these models is closely tied to their ability to reproduce the evidence on MPCs.

As we explain below, there exists a tension between data on the household wealth distribution and the MPCs generated by baseline versions of these models. Our goal in this article is to resolve some of the disagreement and confusion among economists about the sources of this tension and the alternative solutions that have been proposed.

In most of our analysis, we will focus on the quarterly MPC out of a \$500 windfall. This size and frequency align with most of the empirical evidence on MPCs and correspond roughly to the size of fiscal stimulus payment policies. A large body of empirical evidence indicates that the average quarterly MPC on nondurable goods and services out of transitory income changes of \$500–\$1,000 is between 15% and 25% (see Jappelli & Pistaferri 2010 and Havranek & Sokolova 2020 for surveys).¹ Moreover, this average masks substantial heterogeneity in MPCs. Many households have MPCs that are close to zero, but some households have MPCs not far from 1, with a great deal of variation in between (see, e.g., Misra & Surico 2014, Gelman 2021, Lewis et al. 2019). Some of this heterogeneity is explained by the distribution of liquid wealth and some by fixed individual characteristics, but most is left unexplained by observable data. There is also evidence that households respond more strongly both to smaller windfalls than to larger windfalls and to negative income shocks than to positive income shocks.

Despite the importance of MPCs for macroeconomics, until recently this evidence has made only a guest appearance in modern macroeconomics, because of the representative agent assumption implicit in most models. With a representative agent, or complete markets, there is a single common MPC that is roughly the same size as the interest rate, so these models are unable to speak to the evidence. In the last two decades, macroeconomic models have evolved. Today, the heterogeneous agent incomplete-market framework is widely accepted to be a more useful model of the household sector than the representative agent complete-market framework. As a result, MPCs have regained an important role as a point of connection between models and data. Relative to representative agent models, heterogeneous agent models can deliver two realistic features of MPCs: (a) The aggregate MPC can be much larger than the interest rate, and (b) the distribution of MPCs can be very dispersed.

¹Three approaches have been used to measure empirical MPCs: (a) quasi-experimental evidence, exploiting, for example, historical episodes of economic impact payments (Johnson et al. 2006, Parker et al. 2013, Kueng 2018, Aydin 2022), government shutdowns (Gelman et al. 2020), and lottery wins (Fagereng et al. 2021, Golosov et al. 2021); (b) survey instruments that pose hypothetical questions to households about how they would spend income windfalls (Shapiro & Slemrod 2003, 2009; Jappelli & Pistaferri 2014; Bunn et al. 2018; Parker & Souleles 2019; Fuster et al. 2021); and (c) semistructural methods that identify transitory income changes (Blundell et al. 2008, Andersen et al. 2020, Ganong et al. 2020, Commault 2022).

From a theoretical perspective, this class of models can generate large and heterogeneous MPCs due to three characteristics: (a) hand-to-mouth (HtM) households, (b) precautionary saving, and (c) ex-ante heterogeneity. First, these models feature poor HtM households—that is, households with zero wealth and borrowers for whom credit constraints bind. These types of households consume the majority of any extra liquidity they receive. Two-asset versions of the models with liquid and illiquid assets also feature wealthy HtM households that hold the bulk of their wealth in illiquid assets. These households' consumption responds to small windfalls, similarly to the consumption of poor HtM households (Kaplan & Violante 2014, Kaplan et al. 2014).

Second, uninsurable risk, combined with a consumption smoothing motive from concave utility, creates two forces that increase the MPC. If the utility function displays prudence ($u''' > 0$), households save more to avoid the consequences of future negative income shocks (Carroll & Kimball 1996). Even in the absence of prudence, households that face an occasionally binding liquidity constraint will engage in precautionary saving to avoid hitting the constraint (Zeldes 1989, Deaton 1991, Carroll & Kimball 2001). These precautionary saving motives are especially strong at low levels of wealth but disappear as wealth grows sufficiently large. As a result, precautionary motives lower the level of consumption relative to the certainty equivalent case, which is linear in wealth but does so more at low wealth levels. The consumption function is therefore concave in wealth, and MPCs can be large for low-wealth households, even for those that are not HtM. Carroll (2001) and Carroll et al. (2021) provide lucid expositions of these arguments and exhaustive surveys of the literature.

Third, households can differ due to ex-ante characteristics, as opposed to random circumstances. Impatience, high willingness to substitute, low returns on saving, and behavioral biases can all be sources of high MPCs (Carroll et al. 2017, Aguiar et al. 2020). These sources of heterogeneity affect MPCs both through their effect on the wealth distribution (and therefore the fraction of households in the former two circumstances) and through their effect on the consumption function at a given level of wealth.

But how close can this class of models come to matching the evidence on MPCs, when constrained to be consistent with key features of the wealth and earnings distributions? Which features of the models are most important for their quantitative performance? By solving multiple versions of the model under different calibration strategies, we arrive at six main findings, described below.

1. One-asset precautionary saving models calibrated to be consistent with measures of aggregate wealth in US data yield aggregate quarterly MPCs between 3% and 5%. This is an order of magnitude bigger than in corresponding representative agent models but much lower than in their empirical counterparts.
2. Calibrations that target liquid wealth or the share of HtM households can generate an average MPC that is as large as, or larger than, in the data. However, these versions of the model abstract from more than 98% of wealth in the economy and so are of limited use in general equilibrium applications.
3. Extensions of the one-asset model that incorporate either ex-ante heterogeneity or behavioral preferences can generate MPCs as large as in the data, while remaining consistent with the amount of aggregate wealth. However, these versions of the model vastly understate the amount of wealth held by households in the middle of the distribution and feature an excessively polarized wealth distribution. For example, median wealth is 5 to 10 times smaller than in the data. We call this the missing-middle problem.
4. Judiciously chosen versions of spender-saver models can generate MPCs as large as in the data, without the missing-middle problem. However, these models have other implausible implications, such as matching the size and sign dependence of MPCs observed in data.

5. Two-asset precautionary saving models can reconcile all these tensions, but baseline versions require a sizable gap in financial returns between illiquid and liquid assets.
6. Extensions of the two-asset model that introduce nonfinancial benefits from holding illiquid assets can simultaneously match the MPC and wealth data, with smaller differences in financial returns. Examples include direct utility flows, such as the benefits of owner-occupied housing, and indirect benefits, such as those that arise in models of temptation in which households value the commitment that comes from holding wealth in illiquid assets.

The rest of the review is organized as follows. Section 2 analyzes the canonical one-asset incomplete-market model. Section 3 extends the one-asset model in several directions. Section 4 analyzes the two-asset model. Section 5 discusses other concepts of MPCs. Section 6 concludes. The **Supplemental Appendix** contains additional details about data, models, calibrations, and simulations.

2. ONE-ASSET MODELS

In this section, we present findings from various versions of a standard one-asset precautionary saving model.

2.1. Baseline One-Asset Model

We focus mostly on a discrete-time version of the one-asset models, but we also present results for a continuous-time version for consistency with the two-asset models in Section 4.

2.1.1. Environment. The economy is populated by a continuum of households of measure one that survive each period with probability $(1 - \delta)$. Conditional on surviving, a household's time preference factor is given by $\tilde{\beta}$, implying an effective discount factor of $\beta = \tilde{\beta}(1 - \delta) < 1$. Period utility is given by $u(c_t)$, where u is strictly increasing and concave, and c_t denotes consumption expenditures. At each date t , households are endowed with labor income y_t , which follows an exogenous stochastic process described below. Income draws are independent and identically distributed (IID) across households. Households can save, but not borrow, in a risk-free asset b_t with rate of return $R = 1 + r$. The household problem is

$$\max_{\{c_t\}} \mathbb{E}_0 \sum_{t=0}^{\infty} \beta^t u(c_t), \quad 1.$$

subject to

$$c_t + b_{t+1} = Rb_t + y_t, \quad b_{t+1} \geq 0, \quad b_0 = 0.$$

The solution to the household problem yields decision rules for consumption, $c(b, y)$, and next period's wealth, $b'(b, y)$. These decision rules induce a stationary distribution that we denote by $\mu(b, y)$, with associated marginal distribution $\mu(b)$ over wealth.²

2.1.2. Parameterization. Consistent with our focus on quarterly MPCs, our baseline discrete-time model has a period of one quarter. We set $\delta = 1/200$, so that the expected adult life span is

²If the income process has multiple components, then with a slight abuse of notation y should be intended as the vector of all income states.

50 years. In our baseline parameterization we assume a constant elasticity utility function $u(c) = c^{1-\gamma}$, with $\gamma = 1$ so that $u(c) = \log c$. We set $b = 0$ so that there is no borrowing, and we set the interest rate to $r = 0.0025$, or 1% per year. This partial equilibrium approach keeps our exercise especially clear because it allows us to move the interest rate independently of the discount factor and to highlight their respective importance in determining MPCs.³ We model the process for log income as the sum of two orthogonal components, an autoregressive [AR(1)] component and an IID component. We assume that shocks to both components arrive stochastically with a Poisson arrival rate of 1/4, so that shocks are received on average once a year. We estimate the remaining parameters of the income process by matching moments of the household labor income distribution from the Panel Study of Income Dynamics (see **Supplemental Appendix A.1** for details).

2.1.3. Wealth distribution. We choose the effective discount factor β so that mean wealth in the stationary distribution of the model is consistent with mean wealth in the United States. We express all values as multiples of mean annual household earnings, which we define as labor income plus social security income for retired households. Our income and wealth statistics come from the 2019 Survey of Consumer Finances, from which we exclude households in the top 5% of the wealth distribution (see **Supplemental Appendix A.3** for details).⁴ In the bottom 95% of the population, mean annual earnings is \$67,000 and mean net worth is \$275,000, or 4.1 times the annual earnings.

Panel B in **Table 1** reports key wealth statistics from the data (column 1) and in the baseline model (column 2), with the discount factor chosen to hit this target. The implied annualized value for β is 0.98. Median wealth in the model (1.3) is quite close to its value in the data (1.5), despite not being explicitly targeted. The baseline model generates fewer low-wealth households than in the data. A common definition of HtM households in the literature is households whose wealth is less than half their monthly income (Kaplan et al. 2014). According to this definition, only 2% of households in the model are HtM compared with 14% in the data, because optimizing households seek to save away from HtM regions of the asset space (see the “Share of HtM” row in **Table 1**).⁵

2.1.4. Marginal propensities to consume in the baseline model. Our main object of interest is the quarterly MPC out of a one-time unanticipated windfall of size x . For a household with state vector (b, y) at the time when the windfall is received, the impact MPC (or MPC at horizon 0) is

$$m_0(x; b, y) = \frac{c(b + x, y) - c(b, y)}{x}. \quad 2.$$

³In the equilibrium of a closed economy the two would be tightly connected, once a target for the wealth-income ratio is chosen. A number of recent papers in the literature (Auclert et al. 2018, Kaplan et al. 2018, Wolf 2020) have demonstrated that one can usefully separate macro questions about the transmission of shocks and the effects of policies into a partial equilibrium response and a general equilibrium amplification. The analysis in this review is purely about the size of the initial household response to an income shock and about how different model assumptions matter, not about their general equilibrium implications.

⁴The top 5% holds 65% of the total net worth in the economy. We exclude this group from our calibration because it is well understood that the simple precautionary saving models we consider here are not well suited to explain the top tails of the wealth distribution (De Nardi & Fella 2017, Benhabib & Bisin 2018). Moreover, the top of the wealth distribution has a negligible impact on our definition of the MPC.

⁵This feature is robust to other definitions. Fewer than 1% of households have zero or negative wealth, compared with 11% in the data; 3% hold less than \$1,000, compared with 15% in the data; and 12% hold less than \$5,000, compared with 20% in the data.

Table 1 Baseline one-asset model and sensitivity analysis with respect to which moment of the wealth distribution is targeted to set the discount factor

	(1) Data	(2) Baseline	(3) Mean(<i>a</i>)	(4) Median(<i>a</i>)	(5) Mean(<i>a</i>)	(6) Median(<i>a</i>)	(7) HtM
Quarterly MPC (%)	NA	4.6	2.7	4.3	14.0	33.7	22.0
Annual MPC (%)	NA	14.6	8.5	13.6	40.8	77.4	58.7
Quarterly MPC of the HtM (%)	NA	28.7	26.3	28.4	33.2	42.1	36.6
Share of HtM (%)	14.2	2.5	1.7	2.3	7.5	37.4	14.2
Annualized discount factor	NA	0.980	0.988	0.981	0.945	0.826	0.909
Panel A: Decomposition							
Gap with baseline MPC	NA	NA	−1.9	−0.3	9.4	29.1	17.4
Effect of MPC function	NA	NA	−0.8	−0.1	3.1	11.1	6.0
Effect of distribution	NA	NA	−1.3	−0.2	4.6	12.5	7.7
Interaction	NA	NA	0.2	0.0	1.7	5.6	3.7
Panel B: Wealth statistics							
Mean wealth	4.1	4.1	9.4	4.6	0.6	0.1	0.3
Median wealth	1.5	1.3	3.5	1.5	0.2	0.0	0.1
<i>a</i> ≤ \$1,000	15.1	2.5	1.8	2.4	6.2	22.2	10.5
<i>a</i> ≤ \$5,000	19.5	11.6	7.5	10.9	28.6	62.3	42.1
<i>a</i> ≤ \$10,000	24.6	18.5	11.8	17.4	44.3	77.2	60.2
<i>a</i> ≤ \$50,000	37.8	40.3	26.8	38.2	79.3	96.4	91.1
<i>a</i> ≤ \$100,000	49.4	51.9	36.0	49.4	89.8	99.0	97.1
Wealth, top 10% share	49.9	46.6	44.6	46.1	52.0	56.5	51.6

Column 1 contains some key moments of the wealth distribution computed from the 2019 Survey of Consumer Finances (SCF) (<https://www.federalreserve.gov/econres/scfindex.htm>), excluding the top 5%. Column 2 is the baseline model where we set the discount factor to match mean net worth in our SCF sample. Columns 3 and 5 match mean net worth including the top 10% and mean liquid wealth, respectively. Columns 4 and 6 match median net worth and median liquid wealth. Column 7 matches the share of HtM households. The top 10% share in column 1 is computed from our sample, which includes the bottom 95% of households in the 2019 SCF. Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume; NA, not applicable.

Recall that, without income uncertainty or borrowing constraints, the consumption function is linear in wealth, with constant slope m_0^* , as given by the following formula:⁶

$$c_0 = m_0^* \left[Rb_0 + \sum_{t=0}^{\infty} \left(\frac{1}{R} \right)^t y_t \right], \text{ with } m_0^* = 1 - R^{-1} [R\beta]^{\frac{1}{\gamma}}. \tag{3}$$

We refer to m_0^* as the certainty MPC. It is decreasing in the discount factor and, provided that $\beta R < 1$, is decreasing in the intertemporal elasticity of substitution (IES) $1/\gamma$. With log utility $\gamma = 1$, the MPC is equal to the effective discount rate, $m^* = 1 - \beta$, which in our baseline calibration would give a quarterly MPC of 0.5%, independently of the size of the windfall x .

When computing the MPC in our models with uninsurable income risk, we set $x = \$500$ in Equation 2, which is the approximate size of common fiscal stimulus programs from which MPCs were measured in the literature.⁷

Average MPCs are reported in **Table 1**. The average quarterly MPC in the baseline precautionary saving model is 4.6%, nearly one order of magnitude larger than the certainty

⁶Readers are referred to **Supplemental Appendix B.3** for the derivation of Equation 3.

⁷In Section 5 we explore MPCs out of other size windfalls and losses, over other durations.

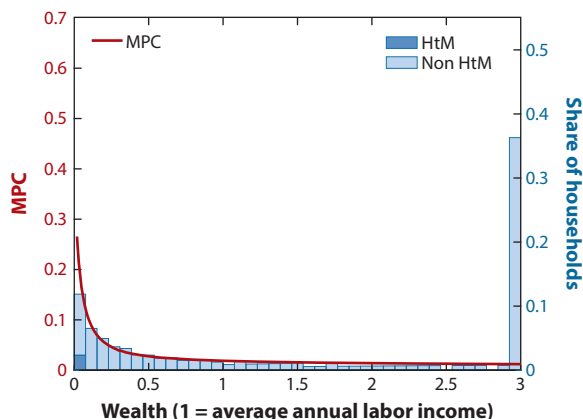


Figure 1

The MPC function (left scale) and the household wealth distribution (right scale) for the baseline one-asset model. Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume.

MPC. **Figure 1** displays the MPC as a function of wealth for a household with mean income, m_0 (\$500; b, \bar{y}), superimposed over the stationary wealth distribution. Note how the MPC quickly converges to the MPC under certainty, $1 - \beta$, as wealth rises. The quarterly MPC of HtM households is very high, nearly 29%, but at a level of wealth of around 0.5 (about \$35,000), the effects of concavity of the consumption function on the MPC have already dissipated: It is only households with very low levels of wealth that contribute to generating an MPC that is substantially above $1 - \beta$. This is a theme that will resurface as we explore richer versions of the model.

2.1.5. Decomposition relative to certainty benchmark. To better understand why the precautionary saving model with uncertainty generates a larger average MPC (4.6% versus 0.5%), we propose a decomposition of the difference between the MPCs in the two models. Let $m_0^{BC}(b)$ be the MPC function in a model that is identical to the baseline model, except that household income is deterministic and all households receive the average level of income. Despite the absence of income risk, if $\beta R < 1$ the consumption function in this model is concave because of the presence of a borrowing constraint (Helpman 1981). The average MPC in the baseline model \bar{m}_0 can then be written as

$$\bar{m}_0 = \underbrace{m_0^*}_{\text{certainty}} + \underbrace{\int_B [m_0^{BC}(b) - m_0^*] d\mu(b)}_{\text{borrowing constraints}} + \underbrace{\int_{B \times \mathcal{Y}} [m_0(b, y) - m_0^{BC}(b)] d\mu(b, y)}_{\text{precautionary savings + income risk}}.$$

The first term is the certainty MPC, the second term captures the role of borrowing constraints in the absence of income risk, paired with the desire to frontload consumption (i.e., $\beta R < 1$), and the final term captures the additional effect of uninsurable risk and precautionary savings.

This decomposition (see panel A in **Table 1**) reveals that the bulk of the gap with the certainty MPC is explained by the presence of borrowing constraints together with a declining optimal consumption profile (68%), even in the absence of idiosyncratic uncertainty. The residual (38%) is accounted for by the precautionary saving motive induced by prudence. Thus, perhaps

surprisingly, a strong force toward concavity of the consumption function emerges even without income uncertainty.⁸

2.2. Alternative Calibrations of the Baseline

We conduct an extensive sensitivity analysis of the baseline model. Unless otherwise specified, we always recalibrate the discount factor so that each version of the model has the same ratio of mean wealth to mean earnings of 4.1.

2.2.1. Deviations between heterogeneous agent models. To understand why different versions of the baseline model yield different MPCs, we propose a decomposition of the difference in average MPCs across models. The average MPC can be larger in one model than in another either because the consumption function is steeper, or because the distribution of households is more concentrated in steeper regions of the state space. Let μ^* and \mathbf{m}_0^* be the distribution and average MPC in the baseline model from Section 2.1, respectively. We can then write the average MPC in an alternative model as

$$\begin{aligned} \bar{m}_0 = & \underbrace{\mathbf{m}_0^*}_{\text{benchmark MPC}} + \underbrace{\int_{B \times \mathcal{Y}} [\mathbf{m}_0(b, y) - \mathbf{m}_0^*] d\mu^*(b, y)}_{\text{consumption function}} + \underbrace{\int_{B \times \mathcal{Y}} \mathbf{m}_0^*(b, y) [d\mu(b, y) - d\mu^*(b, y)]}_{\text{distribution}} \\ & + \underbrace{\int_{B \times \mathcal{Y}} [\mathbf{m}_0(b, y) - \mathbf{m}_0^*] [d\mu(b, y) - d\mu^*(b, y)]}_{\text{interaction}}. \end{aligned}$$

The component labeled “consumption function” captures the difference in average MPCs that arises because the consumption functions are different in the two models. The component labeled “distribution” captures the difference in average MPCs that arises because the stationary distributions of the two models put different mass in different parts of the state space. The component labeled “interaction” arises because these two effects may reinforce or offset each other.

2.2.2. Target for wealth-income ratio. In **Table 1** we report MPCs under alternative aggregate wealth targets. We start by considering a higher target for mean wealth based on the full population without dropping the wealthiest 5% of households, which we set at \$750,000, corresponding to a ratio of wealth to mean earnings of 9.4 (column 3). The average MPC under this calibration drops to 2.7%, driven mostly by the smaller fraction of low-wealth households. With this calibration, median wealth is 3.5, over twice as large as in the data. Next, we target median wealth instead of mean wealth, which gives an average MPC that is almost unchanged from the baseline because median wealth in the baseline model is already close to the value in the data (column 4).

In columns 5 and 6 we report results from what is known in the literature as a liquid wealth calibration. Rather than measuring wealth in the Survey of Consumer Finances (SCF) as total net worth, we only include liquid wealth, which we define as bank accounts and directly held stocks and bonds, net of credit card debt (see **Supplemental Appendix A.3** for details). The logic behind this approach is that liquid wealth is a better measure of those funds that households can readily

⁸Of course, in equilibrium it is income uncertainty that pushes the interest rate below the discount rate.

access to smooth consumption against unexpected income fluctuations. The ratio of mean liquid wealth to mean income is only 0.56. Choosing the discount factor to match this alternative target for mean wealth, the quarterly MPC is 14% (column 5), which is much higher than in the baseline and closer to empirical estimates. When we target median liquid wealth (\$3,100, or 0.05 times the average earnings), the average quarterly MPC rises to 33% (column 6). A related strategy is to choose the discount factor to match directly the fraction of HtM households in the data. When we target 14% of HtM households, the average quarterly MPC is 22%, in between the mean and median liquid wealth calibrations (column 7).

In all three calibrations, the discount factor is much smaller than in the baseline. The lower discount factor raises both the number of households with low wealth around the concave region of the consumption function and the certainty MPC, $(1 - \beta)$ —i.e., the slope of the consumption function for high wealth levels. The MPC decomposition in **Table 1**, panel A confirms this intuition: The steeper consumption function and the shift of the distribution play similar roles in accounting for the higher MPC.

Despite the apparent success of liquid wealth calibrations at generating high MPCs, it is important to note that these calibrations are difficult to integrate into modern dynamic macroeconomic models, because they compel one to abstract from essentially the entire stock of aggregate assets owned by the household sector. This limits the usefulness of these calibrations in general equilibrium models with capital (land, housing, or productive capital). For example, the calibration that matches mean liquid wealth in the data (column 5) effectively abstracts from 85% of the wealth in our baseline sample (which excludes the top 5%), or 98% of total wealth. The calibrations in columns 6 and 7 miss an even larger share of total wealth in the economy.

2.2.3. Interest rate. **Table 2** shows that lowering the interest rate from 1% per annum (p.a.) to 0% p.a. has a negligible effect on the MPC. Raising the interest rate to 5% p.a. increases the average quarterly MPC by around half a percentage point (pp). A higher interest rate leads to a calibration with a lower discount factor, which raises the certainty MPC and makes the consumption function steeper. This is confirmed by the decomposition, which shows that the entirety of the difference in MPC relative to the baseline is due to the consumption function.

2.2.4. Curvature in utility. **Table 2** also shows that changing the curvature parameter of the coefficient of relative risk aversion (RRA) utility function γ away from $\gamma = 1$ has only a small effect on the average MPC. Higher RRA and lower IES (higher γ) strengthen the precautionary saving motive and concavify the consumption function. Because of this higher desire for saving, a lower discount factor is required to generate the same amount of aggregate wealth. These forces increase the MPC at all wealth levels. There is, however, an offsetting force: The stronger precautionary motive means that there are fewer households close to the borrowing constraint in the stationary distribution for a given aggregate amount of wealth. Column 5 shows that for a coefficient of RRA of $\gamma = 6$, this latter effect dominates, and the average quarterly MPC falls from 4.6% to 3.0%. The decomposition shows that the 1.6-pp lower MPC is comprised of a 0.5-pp point increase from the lower discount factor, offset by a 2.1-pp fall from the smaller fraction of low-wealth households. For example, relative to the baseline where $\gamma = 1$, median wealth is almost twice as large in the $\gamma = 6$ economy. The assumption of a constant RRA utility function means that this analysis cannot separate the effects of changing the RRA from those of changing the IES. In Section 3 we generalize preferences and perform this analysis.

2.2.5. Model frequency. In **Supplemental Table E.1** we also report results for the baseline calibration with a model period of one year rather than one quarter, and for a continuous-time

Table 2 Baseline one-asset model and sensitivity analysis with respect to the interest rate r (annualized values) and the RRA coefficient γ in the utility function

	(1) Baseline	(2) $r = 0\%$	(3) $r = 5\%$	(4) RRA = 0.5	(5) RRA = 6
Quarterly MPC (%)	4.6	4.4	5.2	5.3	3.0
Annual MPC (%)	14.6	13.9	16.8	16.4	10.6
Quarterly MPC of the HtM (%)	28.7	28.5	29.5	29.0	26.6
Share of HtM (%)	2.5	2.4	2.5	3.3	0.6
Annualized discount factor	0.980	0.990	0.944	0.986	0.840
Panel A: Decomposition					
Gap with baseline MPC	NA	−0.2	0.6	0.7	−1.6
Effect of MPC function	NA	−0.2	0.7	0.1	0.5
Effect of distribution	NA	0.0	−0.1	0.6	−2.1
Interaction	NA	0.0	0.0	0.0	0.1
Panel B: Wealth statistics					
Mean wealth	4.1	4.1	4.1	4.1	4.1
Median wealth	1.3	1.4	1.4	1.1	2.3
$a \leq \$1,000$	2.5	2.5	2.4	3.4	0.6
$a \leq \$5,000$	11.6	11.6	11.1	13.9	3.1
$a \leq \$10,000$	18.5	18.4	17.8	21.4	6.7
$a \leq \$50,000$	40.3	40.1	39.5	43.3	24.8
$a \leq \$100,000$	51.9	51.6	51.2	54.2	38.8
Wealth, top 10% share	46.6	46.2	45.9	48.6	36.3

Column 1 is the baseline model with interest rate r equal to 1% and coefficient of RRA equal to 1. In columns 2 and 3 we solve the model with different values of the interest rate. In columns 4 and 5 we solve the model with different values of the coefficient of RRA. Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume; NA, not applicable; RRA, relative risk aversion.

version of the model. The average annual MPC in the annual model is 14.3%, essentially the same as the average annual MPC in the quarterly model, which is 14.6%. Note that this value is lower than what one would obtain by cumulating the quarterly MPC over four quarters, using the formula $m_4 = (1 + m_0)^4 - 1$, which is often erroneously used in the literature. Applying this formula would yield an annual MPC of 19.7%, or 35% higher than the correct annual MPC. Simply multiplying the quarterly MPC by four would give an annual MPC of 18.4%, or 26% higher than the correct MPC.

The quarterly MPC in the continuous-time version of the model is 3.0%, slightly lower than our baseline discrete-time model.⁹

2.2.6. Income process. We also considered several alternative income processes to the one in our baseline model, including: (a) assuming that income shocks arrive on average once each

⁹Discrete- and continuous-time models necessitate different specifications of the income process. We have adopted analogous processes calibrated to match the same targets, so as to enable comparisons between the two. **Supplemental Table A.2** contains parameter estimates for the continuous-time counterpart of the income processes. Wealth and MPC statistics are very similar, which ensures that, when we move to the continuous-time two-asset model in Section 4, the differences are not being driven by the switch to continuous time.

quarter, rather than once each year; (b) estimating the arrival rates of the shocks alongside the other parameters by targeting the kurtosis of income growth rates at different lags, in addition to the variance of income growth rates; (c) estimating an annual income process and converting the parameter estimates into quarterly values using the approach of Krueger et al. (2016); and (d) eliminating transitory shocks. We also considered alternative processes in the annual version of the model, including a version with higher persistence in the AR(1) component and a version with individual-specific fixed effects. Details for all these statistical representations of income dynamics can be found in **Supplemental Table A.2** and a summary of the results can be found in **Supplemental Table E.2**.

For most of these alternative specifications of the income process, the MPCs are extremely close to the baseline model. The only versions that differ significantly from the baseline are those in which we assume no transitory shocks and those in which we assume a high persistence in the AR1 component. Without transitory shocks, households are willing to hold less precautionary wealth because they face less easily insurable income risk. As a result, this calibration has a share of HtM consumers of 18% compared to 2% in the baseline and a larger average MPC. In the second specification, the AR1 component is very close to a random walk. As a result, the estimation loads all the short-term volatility onto a higher transitory variance. Both forces contribute toward a lower MPC: Because shocks are near-permanent, a higher discount factor is needed to match aggregate wealth; and because of the sizable transitory variance, the distribution of wealth has less mass near the credit limit.

2.2.7. Survival and bequests. We also examined sensitivity to different assumptions about the survival rate, how the assets of the deceased are distributed, and how the assets of new-born households are determined. As long as the discount factor is always recalibrated to match the same amount of aggregate wealth, none of these assumptions matters (see **Supplemental Table E.3**).

3. EXTENSIONS OF THE ONE-ASSET MODEL

3.1. Ex-Ante Heterogeneity

In this section, we extend the one-asset model to allow for various forms of ex-ante heterogeneity.

3.1.1. Heterogeneity in discount factors. We start by allowing for heterogeneity in households' discount factors β . We consider a discretized uniform distribution for β (annualized) with five equally spaced grid points between $[\beta - 2\Delta, \beta + 2\Delta]$. We choose the the midpoint β to match average wealth of 4.1, as in our previous calibrations. **Table 3** reports results for versions with a moderate amount of heterogeneity ($\Delta = 0.02$, column 2) and with a large amount of heterogeneity ($\Delta = 0.04$, column 3). For both calibrations, the median effective discount factor is substantially lower than in the baseline economy. The reason is that in these economies there is a subset of very patient households that have a strong intertemporal saving motive, and it is these households that hold the bulk of aggregate wealth. This can be seen from the much higher share of wealth held by the top 10% households (88% and 79%), compared to the model without discount factor heterogeneity (47%). The model can therefore match the target for mean wealth while still allowing for a large fraction of households to be impatient. Consequently, the models with discount factor heterogeneity match well the very bottom of the wealth distribution. For example, in the calibration with $\Delta = 0.04$, 14% of households are HtM, as in the data.

The average quarterly MPC in the models with discount factor heterogeneity are much higher than in the baseline model, and with enough heterogeneity, these models can approach the target empirical values. With $\Delta = 0.04$, the average MPC is nearly 19%, four times as large as in the baseline model. Whereas both the shape of the MPC function and the stationary wealth

Table 3 One-asset model with heterogeneity in the discount factor β (annualized values) and in the rate of return r (annualized values)

	(1) Baseline	(2) Het β	(3) Het β	(4) Het β	(5) Het β	(6) Het r	(7) Het r	(8) S-S
Set of β	NA	0.91,...,0.99	0.83,...,0.99	0.85,...,1.01	0.88,...,1.04	NA	NA	NA
Switch probability β	NA	0.00	0.00	0.02	0.10	NA	NA	NA
Set of r	NA	NA	NA	NA	NA	-1, 1, 3	-3, 1, 5	NA
Quarterly MPC (%)	4.6	11.7	18.6	14.3	6.7	7.5	11.6	16.8
Annual MPC (%)	14.6	33.4	47.8	38.1	19.6	22.8	33.1	26.1
Quarterly MPC of the HtM (%)	28.7	33.8	38.5	37.0	33.0	31.0	33.8	85.4
Share of HtM (%)	2.5	6.7	14.2	9.9	4.1	4.0	6.8	14.8
Median annualized β	0.980	0.953	0.915	0.933	0.963	0.970	0.954	0.392
Panel A: Decomposition								
Gap with baseline MPC	NA	7.1	14.0	9.7	2.1	2.9	7.0	12.2
Effect of MPC function	NA	2.4	4.3	3.4	1.2	1.2	2.7	1.9
Effect of distribution	NA	3.2	6.0	4.1	0.8	1.4	3.1	3.3
Interaction	NA	1.4	3.7	2.1	0.1	0.3	1.2	7.0
Panel B: Wealth statistics								
Mean wealth	4.1	4.1	4.1	4.1	4.1	4.1	4.1	4.1
Median wealth	1.3	0.3	0.2	0.3	1.1	0.6	0.3	1.0
$a \leq \$1,000$	2.5	5.5	10.0	7.6	3.8	3.7	5.5	14.6
$a \leq \$5,000$	11.6	24.1	35.9	28.2	14.9	16.8	23.7	22.8
$a \leq \$10,000$	18.5	36.5	49.7	39.9	22.2	26.6	35.8	28.7
$a \leq \$50,000$	40.3	64.4	74.5	64.1	43.5	52.6	62.5	46.6
$a \leq \$100,000$	51.9	73.8	81.0	72.2	54.6	63.6	71.5	56.2
Wealth, top 10% share	46.6	79.2	87.9	75.7	51.4	64.7	74.0	50.9

Column 1 is the baseline model with common discount factor and interest rates. Columns 2 and 3 are models with fixed ex-ante heterogeneity in discount factors across households. Columns 4 and 5 solve models in which the individual discount factor switches randomly. Columns 6 and 7 are models with fixed ex-ante heterogeneity in interest rates across households. Column 8 is the spender-saver model. Abbreviations: Het, heterogeneity; HtM, hand-to-mouth; MPC, marginal propensity to consume; NA, not applicable; S-S, spender-saver.

distribution contribute to the higher MPC, the majority of the effect comes from the larger fraction of low-wealth households (see the decomposition in panel A). **Figure 2** displays the MPC functions for high- and low- β households overlaid on the stationary wealth distribution for the two groups of households. The figure illustrates how the impatient households not only are amassed near to the borrowing constraint but also have higher MPCs at all levels of wealth.

We also examined a version of the model in which households switch randomly between different discount factors. We assumed that a household draws a new value of β with probability p , independent of its current value.¹⁰ The models with stochastic discount factors generate a lower average MPC compared with a model with the same stationary distribution of β but with fixed heterogeneity. The reason is that in the model with stochastic discount factors there is a weaker correlation between wealth and discount factors. In the stochastic β model, some low-wealth

¹⁰We set $\Delta = 0.01$ and consider annual switching rates of 0.02, so that the expected duration of a discounting regime has the same expected duration of a lifetime (column 4), and of 0.1., so that the expected duration is equal to a decade (column 5).




Figure 2

The MPC function (left scale) and the household wealth distribution (right scale) for the model with large heterogeneity in discount factors across households. Abbreviations: β , discount factor; MPC, marginal propensity to consume.

households that were previously impatient become patient and quickly accumulate wealth. The wealth distribution is therefore less concentrated at the bottom, which is evident in the lower share of HtM households and higher value of median wealth.

Although the models with discount factor heterogeneity feature a large average MPC while generating the key features of the two tails of the wealth distribution, these models typically fail in reproducing the wealth distribution everywhere in between. For example, in the model with $\Delta = 0.04$, despite matching the fraction of wealth-poor households (14% HtM, 10% with wealth lower than \$1,000), the wealth distribution is excessively squeezed toward the bottom and has far too many households with wealth just above this threshold: Median wealth is 10 times smaller than in the data, and 75% of households have wealth below \$50,000 compared with only 38% in the data. This problem, which we label the missing middle, is a recurring issue that arises in many of the versions of the model with heterogeneity that we examine below. Comparing the wealth distributions in **Figures 1 and 2** clearly illustrates this shortcoming.¹¹

3.1.2. Heterogeneity in rates of return. In columns 6 and 7 of **Table 3**, we report results for an economy with fixed heterogeneity in r , uniformly distributed over $\{-1\%, 1\%, 3\%\}$ p.a. (column 6) and over $\{-3\%, 1\%, 5\%\}$ p.a. (column 7). Heterogeneity in rates of return generates similar results to heterogeneity in discount factors. For example, in the calibration in column 7, 7% of households are HtM, the top 10% share is 74%, and the aggregate quarterly MPC reaches nearly 12%. These economies, however, also feature the missing-middle problem, with median wealth levels much lower than in the data.

3.1.3. Heterogeneity in risk aversion and intertemporal elasticity of substitution. We now allow for heterogeneity in the curvature parameter γ in the coefficient of RRA utility function. The distribution of γ is a discretized uniform with five geometrically spaced grid points in the interval $[\frac{1}{\bar{\gamma}}, \bar{\gamma}]$. In **Supplemental Table E.4**, we report results for a moderate amount of

Supplemental Material >

¹¹An interesting source of ex-ante heterogeneity, which we have not explored because of space limitations, is the income process. Because a model in which households face different expected growth rates in their labor income is isomorphic to one with heterogeneous discount factors, our conclusions extend to the former case.

heterogeneity ($\bar{\gamma} = e^2 = 7.4$) and a large amount of heterogeneity ($\bar{\gamma} = e^3 = 20.1$). In both cases we set the midpoint to $\gamma = e^0 = 1$, as in the baseline model. With enough heterogeneity in γ , these calibrations yield very large average MPCs, because of the high shares of households in the tails of the wealth distribution. For example, a model with values of γ ranging from 0.14 to 7.4 gives an average quarterly MPC of 17%. Most of the wealth is held by the high-RRA/low-IES households, and the top 10% wealth share is around 67%. The large MPC is driven by the low-RRA/high-IES households, and 20% of households are HtM. The decomposition in panel A reveals that it is indeed this different distribution of households relative to the baseline model, rather than the shape of the MPC function, that is the most important factor in generating a larger MPC. But like the previous calibrations with heterogeneity, this economy suffers from the same missing-middle problem: A disproportionate share of households is wealth-poor, and median wealth is much lower than in the data.

To distinguish the roles of risk aversion and intertemporal substitution in generating the high degree of wealth concentration and large MPC in this last calibration, we generalize our preferences to those of Epstein & Zin (1991):

$$U_t = \left\{ (1 - \beta) c_t^{1-\theta} + \beta \left(\mathbb{E}_t \left[U_{t+1}^{\frac{1-\theta}{1-\gamma}} \right] \right)^{1-\gamma} \right\}^{\frac{1}{1-\theta}}, \quad 4.$$

where γ is the coefficient of RRA and $1/\theta$ is the IES. In **Supplemental Table E.5** we report results for various values of γ and θ . Without preference heterogeneity, varying either of these two parameters has only a very small impact on the average MPC once the discount factor is recalibrated to match the same average wealth target.

Allowing for heterogeneity in RRA, holding the IES constant, has barely any effect on the average MPC. However, with heterogeneity in the IES, the model is able to generate large average MPCs. These results are reported in **Table 4** and imply that the findings with a constant RRA utility function described above are being driven by the heterogeneity in the IES rather than in the RRA. For example, with a geometric uniform distribution of θ ranging from $e^{-3} = 0.05$ to $e^3 = 20$, the average quarterly MPC is above 20%. Households with a high IES are willing to absorb consumption fluctuations, and so they hold only small buffer stocks of wealth. However, this calibration suffers from a missing-middle problem, like other versions of the model with ex-ante heterogeneity.

3.2. Behavioral Preferences

In this section we explore two alternative behavioral models for preferences that have been proposed in the literature as ways to generate higher MPCs: temptation and present bias. Our main message is that both models are effectively models of endogenous discount rates, with the discount rate depending on the level of assets. As a result, these models have the same strengths and weaknesses in terms of average MPCs and wealth distribution as the model with ex-ante discount factor heterogeneity discussed in Section 3.1.

3.2.1. Temptation and self-control. Gul & Pesendorfer (2001) propose a model of temptation and self-control in which consumers are tempted to consume according to a preference specification that overweighs current consumption, but they can exert some degree of self-control. These preferences have recently gained popularity in the quantitative macro literature (Krusell et al. 2010, 2002; Nakajima 2017; Pavoni & Yazici 2017; Attanasio et al. 2020). We consider the limiting formulation in which consumers are tempted to consume according to preferences that place

Table 4 One-asset model with heterogeneous Epstein-Zin preferences and one-asset model with Gul-Pesendorfer preferences

	(1) Baseline	(2) Het IES	(3) Het IES	(4) Het IES	(5) Temptation	(6) Temptation	(7) Het temptation
Set of IES $1/\theta$	NA	e^{-1}, \dots, e^1	e^{-2}, \dots, e^2	e^{-3}, \dots, e^3	NA	NA	NA
Set of temptation φ	NA	NA	NA	NA	0.01	0.05	0, 0.05, 0.1
Quarterly MPC (%)	4.6	6.6	10.9	20.8	6.4	19.1	21.4
Annual MPC (%)	14.6	20.1	29.9	39.2	19.0	46.2	47.7
Quarterly MPC of the HtM (%)	28.7	30.5	34.4	66.2	31.4	39.3	45.1
Share of HtM (%)	2.5	4.0	11.1	20.9	4.0	20.4	29.3
Annualized discount factor	0.980	0.977	0.975	0.974	0.982	0.990	0.988
Panel A: Decomposition							
Gap with baseline MPC	NA	2.0	6.3	16.2	1.8	14.5	16.8
Effect of MPC function	NA	0.9	2.1	3.1	0.4	2.7	3.9
Effect of distribution	NA	1.0	3.3	6.0	1.2	7.4	8.3
Interaction	NA	0.2	0.9	7.1	0.2	4.4	4.5
Panel B: Wealth statistics							
Mean wealth	4.1	4.1	4.1	4.1	4.1	4.1	4.1
Median wealth	1.3	0.9	0.5	0.3	1.0	0.1	0.1
$a \leq \$1,000$	2.5	3.7	8.0	19.0	3.9	14.9	22.2
$a \leq \$5,000$	11.6	15.4	24.3	33.3	16.7	42.5	42.8
$a \leq \$10,000$	18.5	23.8	33.7	41.7	25.0	52.6	50.8
$a \leq \$50,000$	40.3	47.3	55.9	60.9	46.2	69.4	66.3
$a \leq \$100,000$	51.9	58.1	64.7	68.3	55.9	73.6	71.5
Wealth, top 10% share	46.6	55.7	62.8	67.9	49.8	70.6	70.1

Column 1 is the baseline model with common IES across households. Columns 2–4 are models with ex-ante heterogeneity in IES. Columns 5 and 6 are models with Gul-Pesendorfer preferences and common temptation parameter. Column 7 is a model with Gul-Pesendorfer preferences and heterogeneous temptation parameter across households. Abbreviations: Het, heterogeneity; HtM, hand-to-mouth; IES, intertemporal elasticity of substitution; MPC, marginal propensity to consume; NA, not applicable.

no weight on future consumption. The degree of temptation is governed by an additional parameter $\varphi \in [0, \infty]$. It is straightforward to show that for $\varphi > 0$, the temptation and self-control model generates a modified Euler equation in which the effective discount factor is decreasing in the average propensity to consume, which has the effect of making poor households act in a more myopic way than wealthier households. For example, with log utility, the effective discount factor becomes

$$\beta \mathbb{E} \left[1 - \frac{\varphi}{1 + \varphi} \left(\frac{c'}{x'} \right) \right],$$

where x is cash on hand. We describe the setup formally in **Supplemental Appendix C.1**.

Columns 5 to 7 in **Table 4** report results for values of the temptation parameter $\varphi \in \{0.01, 0.05\}$. The calibrations with a high degree of temptation can generate a large MPC without any ex-ante heterogeneity, while still matching the target for mean wealth. For example, with $\varphi = 0.05$, the average quarterly MPC is 19%. The decomposition reveals that the shape of the consumption function plays only a minor role, and that the large MPC is due to a different shape of the wealth distribution compared to the baseline model: Low-wealth households have a lower effective discount rate than high-wealth households, which leads to a large mass near the borrowing constraint. In this calibration, median wealth is around \$10,000, i.e., less than one-tenth

Supplemental Material >

Supplemental Material >

of that in the data. Allowing for heterogeneity in φ , centered around a median value of 0.05, has very little effect on either the average MPC or the wealth distribution. The one-asset model with temptation thus also suffers from the missing-middle problem.

3.2.2. Present bias. Starting with Laibson (1997), the other commonly adopted departure from the consumption saving problem with standard preferences is to assume some form of present bias, such as quasi-hyperbolic discounting. We adopt the continuous-time formulation of Laibson et al. (2021), known as instantaneous gratification. Unlike in the model of temptation and self-control, in this case preferences are not time consistent, and so one needs to make an assumption about whether consumers are naive, meaning that they are unaware that their future selves will have different preferences, or are sophisticated, meaning that they are aware of the time inconsistency and play a game against their future selves. Following Laibson et al. (2021), and to give the model its best shot at generating a large average MPC, we assume that households are naive. We describe the setup formally in **Supplemental Appendix C.3**.

The instantaneous gratification model features one additional parameter, $\zeta \leq 1$, which measures the extent of present bias. In **Supplemental Table E.6** we report results for values of $\zeta \in \{0.9, 0.8, 0.7\}$, alongside the baseline continuous-time one-asset model (corresponding to the special case of $\zeta = 1$). The results suggest that once the instantaneous gratification model is recalibrated to match the same level of aggregate wealth, present bias has a negligible effect on the average MPC. In fact, with $\zeta = 0.7$, the average MPC is lower than in the model with exponential discounting. The MPC decomposition reveals that there are two offsetting effects. On the one hand, the instantaneous gratification model has a stationary distribution with a higher number of very low-wealth households (e.g., 19% of households with less than \$1,000, compared with 2.5% in the baseline model). On the other hand, in order to match the same level of aggregate wealth, the calibrated value of the effective discount factor is much larger in the model with present bias (0.997 p.a. versus 0.985 p.a.). This means that everywhere above the borrowing constraint, the MPC is lower in the model with present bias. **Figure 3** illustrates these differences. The decomposition shows that these two effects roughly cancel out.

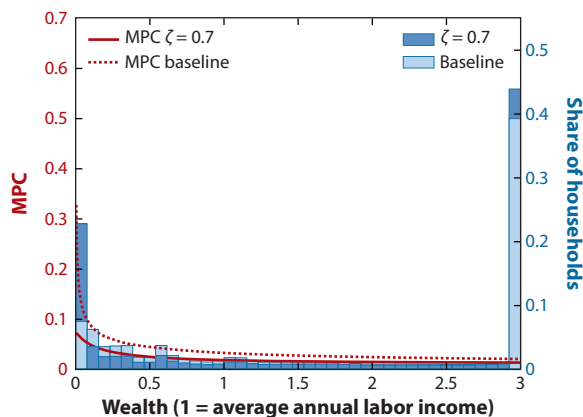


Figure 3

The MPC function (left scale) and the household wealth distribution (right scale) for the model with preferences featuring present bias. Abbreviations: ζ , extent of present bias (no bias if $\zeta = 1$); MPC, marginal propensity to consume.

3.3. Taking Stock

When the one-asset model is calibrated to match the amount of wealth held by the poorest 95% of US households, it generates an average quarterly MPC of around 4%. This value is an order of magnitude larger than the MPC of the representative agent model, but it is still much smaller than empirical estimates.

Without preference heterogeneity or behavioral preferences, the baseline model can only generate an average quarterly MPC of around 20% when it is parameterized to match an amount of aggregate wealth that is 10 to 20 times smaller than its empirical counterpart for the United States, for example in calibrations based on liquid wealth, rather than total wealth. However, these calibrations necessarily ignore almost the entire aggregate capital stock, and so they are of limited use for general equilibrium applications.

When the one-asset model is extended to allow for ex-ante heterogeneity in discount factors, rates of return, or IES across households, it can generate quarterly MPCs of 20%, or higher in some calibrations. In all these versions of the model, the reason for the larger MPC is that the presence of different household types makes it possible for the stationary distributions to contain a substantial number of HtM households, while still being consistent with an average level of wealth as high as in the data. Yet, all these models have far too many households that are not quite poor enough to be HtM but still have very little wealth—the missing middle. For example, the fraction of households with less than \$50,000 is typically much higher than in the data, and the wealth of the median household is 5 to 10 times smaller than in the data.

When preferences are generalized to allow for temptation, the model behaves like one in which discount factors are decreasing in wealth. This feature allows it to generate a high average MPC, but the model suffers from the same missing-middle problem that plagues versions with ex-ante heterogeneity. Versions of the model with present bias introduce offsetting forces when calibrated to match aggregate net worth, and the average MPC remains small.

3.4. Spender-Saver Models

The challenge for one-asset models is therefore to simultaneously generate a sufficiently large amount of wealth in the aggregate and a large enough fraction of HtM households, while still generating enough households in the middle of the distribution. There is one form of ex-ante heterogeneity in the one-asset model that succeeds in this respect. Inspired by the spender-saver model of Campbell & Mankiw (1989), these are versions of the model with discount factor heterogeneity in which there are two groups of households: one with a very low discount factor (the spenders) and one with a high discount factor (the savers). Judiciously chosen calibrations of these models can match all of these features of the data.

To illustrate this, column 8 in **Table 3** reports results from a calibration in which 15% of households have a discount factor of 0.4, and the remaining 85% of households have a discount factor that is chosen so that the ratio of average wealth to average income is 4.1, as in the data. In this spender-saver model, 14.8% of households are HtM, and the average quarterly MPC is 17%. Moreover, median wealth is 0.96, which is not too far from its empirical counterpart of 1.54, and around 47% of households have less than \$50,000 of wealth, compared with 39% in the data. The missing-middle problem is therefore nowhere near as extreme as in other versions of the one-asset model with ex-ante heterogeneity.

Nonetheless, despite its success at generating a large average MPC and simultaneously matching the median and the mean of the wealth distribution, this extreme form of heterogeneity has a major drawback that limits its usefulness for counterfactual policy analysis. For example, fiscal transfers targeted to the impatient group are extremely powerful at triggering an immediate

response of aggregate spending. Because of this feature, the spender-saver model displays an intertemporal MPC function with a steep peak coinciding with the time of the disbursement and little subsequent propagation. We return to these comparisons in Section 5.

4. TWO-ASSET MODELS

In this section, we extend the precautionary saving model to include two assets: a liquid asset with low return, and an illiquid asset with higher return but subject to adjustment frictions. We demonstrate that the two-asset model can go a long way toward resolving the intrinsic tension in the one-asset model, even in the absence of ex-ante heterogeneity.

In two-asset heterogeneous agent models, households can separate their different saving motives into distinct assets. Precautionary and smoothing motives against small, regular income fluctuations induce households to accumulate a buffer of liquid assets, but because the return on this asset is low, a strong negative intertemporal motive ($\beta R \ll 1$) prevents households from accumulating large amounts of liquid wealth. The bulk of household savings is done in the high-return illiquid asset, motivated both by a positive intertemporal motive and by precautionary and smoothing motives against large, infrequent income fluctuations. At endogenous intervals, households move funds between their liquid and illiquid accounts as desired. For instance, in response to a large negative income shock, a household that has exhausted its buffer of liquid wealth can pay a fee or exert effort to withdraw funds from its illiquid account. Similarly, a household that has experienced a long stream of positive income growth and has accumulated excess liquid wealth can transfer funds to its illiquid account, which pays a higher return and is therefore a better vehicle for long-run saving.¹²

Since most households end up holding the majority of their wealth in illiquid assets, which are of limited use for short-term consumption smoothing, they expose themselves to potentially larger consumption fluctuations than if they saved only in liquid assets. However, the welfare cost of these fluctuations is second-order, relative to the first-order gain from earning a higher return on illiquid savings. Because of this trade-off, the model generates wealthy HtM households that have positive, and sometimes substantial, holdings of illiquid wealth, but only small amounts of liquid wealth. Wealthy HtM households coexist alongside the traditional poor HtM households, that hold very little wealth, either liquid or illiquid. Adopting a definition of HtM households analogous to the one used in previous sections (i.e., households with liquid wealth less than half of monthly income), 27% of households were wealthy HtM in the 2019 SCF, in addition to the 14% of poor HtM households. Therefore, in total 41% of US households are HtM. Since wealthy HtM households also have a high MPC out of small one-time windfalls, it is the presence of this additional group of HtM households that enables the two-asset model to generate a substantial average MPC while remaining consistent with the distributions of liquid and illiquid wealth in the data.

4.1. Baseline Two-Asset Model

We formulate the two-asset model in continuous time.

4.1.1. Environment. The economy is populated by a continuum of households that discount the future at the effective rate $\rho = \bar{\rho} + \delta < 1$, where $\bar{\rho}$ is the rate of time preference and δ is

¹²The first scenario corresponds, for example, to a household that suffers a job loss and extracts some equity from its house or withdraws from its retirement account to smooth expenditures. The second scenario corresponds to a household that saves enough to make a down payment on a house.

the death rate. Flow utility is given by $u(c_t)$, where u is strictly increasing and concave, and c_t denotes consumption expenditures. At each date t , households are endowed with log labor income $y_t \in Y$, which follows an exogenous stochastic process described in **Supplemental Appendix B.3**. Households can save, but not borrow, in two assets: a liquid asset b with return r^b , and an illiquid asset a with return $r^a > r^b$. At rate χ , households receive an opportunity to rebalance their financial portfolio, which they can take by paying a fixed transaction cost κ . Between rebalance dates, the illiquid asset accumulates in the background and the household solves a standard consumption saving problem out of its liquid assets. The Hamilton–Jacobi–Bellman (HJB) equation for the household problem is

$$\begin{aligned} \rho v(a, b, y) = & \max_c u(c) + v_b(a, b, y) \dot{b} + v_a(a, b, y) \dot{a} + \mathcal{A}v(a, b, y) \\ & + \chi [v^*(a, b, y) - v(a, b, y)] \end{aligned}$$

subject to

$$\begin{aligned} \dot{b} &= r^b b + y - c, \quad b \geq 0, \\ \dot{a} &= r^a a, \quad a \geq 0, \end{aligned}$$

where \mathcal{A} is the infinitesimal generator of the stochastic process for income. The last term in the HJB equation describes the gain from rebalancing the asset portfolio. The value function after the rebalancing opportunity arises, $v^*(a, b, y)$, is defined by

$$v^*(a, b, y) = \max \{ \omega(a + b, y), v(a, b, y) \}.$$

Upon receipt of an adjustment opportunity, a household will choose to rebalance its portfolio if the gains from rebalancing exceed the transaction cost κ . In this case, the household can choose any feasible combination of liquid and illiquid assets (b', a') , and its value is given by

$$\omega(a + b, y) = \max_{a', b'} v(a', b', y)$$

subject to

$$a' + b' \leq a + b - \kappa, \quad a', b' \geq 0.$$

The solution to this problem yields decision rules for consumption $c(a, b, y)$, portfolio rebalancing $p(a, b, y) \in \{0, 1\}$, the optimal rebalanced portfolio $a'(a, b, y)$ and $b'(a, b, y)$, and the stationary distribution of households $\mu(a, b, y)$.

4.1.2. Parameterization. Relative to the continuous-time one-asset model, there are three additional parameters: the arrival rate for rebalancing opportunities χ , the illiquid asset return r^a , and the transaction cost κ . In our baseline model we set $\chi = 3$ so that households get an opportunity to rebalance on average once per month (the model period is one quarter). We also set the (annualized) liquid return $r^b = -2\%$ (roughly corresponding to a zero nominal return). We then choose the discount rate ρ , the illiquid rate r^a , and the transaction cost κ to match three targets: (a) a ratio of mean total net worth to mean earnings of 4.1, as in the one-asset model; (b) a total share of HtM households (both wealth and poor) of 41%; and (c) a share of poor HtM households of 14%.

Table 5 reports the calibrated parameters and wealth statistics in our baseline two-asset model. The model matches well all three targets. To achieve this, the calibration requires a sizable gap

Table 5 Baseline two-asset model and sensitivity with respect to the liquid and illiquid rates of return

	(1) Baseline one-asset	(2) Baseline two-asset	(3) $r_b = -3\%$	(4) $r_b = 0\%$	(5) $r_a = 5\%$	(6) $r_a = 7.5\%$
Rebalance arrival rate	NA	3.0	3.0	3.0	3.0	3.0
Quarterly MPC (%)	3.0	16.1	17.7	12.3	14.0	16.5
Annual MPC (%)	11.5	41.2	43.7	33.9	37.3	40.9
Quarterly PHtM MPC (%)	33.7	24.3	25.4	27.8	25.8	25.3
Quarterly WHtM MPC (%)	NA	29.8	31.4	22.2	25.6	27.5
Mean MPC at mean wealth (%)	1.7	7.0	7.3	5.9	5.8	7.1
Prob. HtM status at year t and year $t + 1$	0.5	0.9	0.9	0.8	0.8	0.9
Panel A: Calibrated variables						
Annualized discount factor	0.985	0.937	0.937	0.937	0.948	0.926
Annualized liquid return	1.0	-2.0	-3.0	0.0	-2.0	-2.0
Annualized illiquid return	NA	6.2	6.2	6.2	5.0	7.5
Rebalance cost (\$)	NA	1,456.8	1,456.8	1,456.8	1,456.8	1,456.8
Panel B: Targeted statistics						
Mean total wealth	4.1	4.1	4.1	4.1	4.1	4.1
Share HtM	NA	39.5	43.1	28.7	34.6	41.6
Share PHtM	2.0	13.2	15.2	8.5	10.2	15.7
Panel C: Decomposition						
Gap with baseline MPC	NA	NA	1.6	-3.9	-2.1	0.4
Effect of MPC function	NA	NA	0.5	-3.0	-1.3	-0.8
Distributional effect	NA	NA	1.0	-2.7	-1.3	0.2
Interaction	NA	NA	0.1	1.8	0.5	0.9
Panel D: Wealth statistics						
Mean liquid wealth	NA	0.2	0.2	0.3	0.2	0.2
Median total wealth	1.6	1.1	1.1	1.1	1.1	0.9
Median liquid wealth	NA	0.1	0.0	0.1	0.1	0.1
$b \leq \$1,000$	NA	33.0	36.6	23.7	28.7	35.2
$b \leq \$5,000$	NA	54.7	57.0	47.7	52.5	55.3
$b \leq \$10,000$	NA	66.1	67.8	60.7	65.0	66.4
$w \leq \$1,000$	2.5	10.0	11.7	7.1	8.3	12.0
$w \leq \$5,000$	8.0	20.8	21.5	19.3	19.7	21.8
$w \leq \$10,000$	13.4	23.5	23.7	23.3	23.2	24.1
$w \leq \$50,000$	34.9	45.9	46.1	45.2	45.3	47.5
$w \leq \$100,000$	48.0	57.7	57.9	57.6	57.5	59.7
Wealth, top 10% share	40.3	62.0	62.4	60.9	60.9	64.0
Wealth, top 1% share	5.6	19.2	19.6	18.0	17.7	21.0
Gini coefficient, total wealth	0.6	0.8	0.8	0.8	0.8	0.8

Column 1 is the baseline one-asset model. Column 2 is the baseline two-asset model. Columns 3 and 4 are two-asset models solved with a different rate of return on the liquid asset. Columns 5 and 6 are two-asset models solved with a different rate of return on the illiquid asset. b denotes liquid wealth and $w = a + b$ total wealth (or net worth). Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume; NA, not applicable; PHtM, poor hand-to-mouth; Prob., probability; WHtM, wealthy hand-to-mouth.

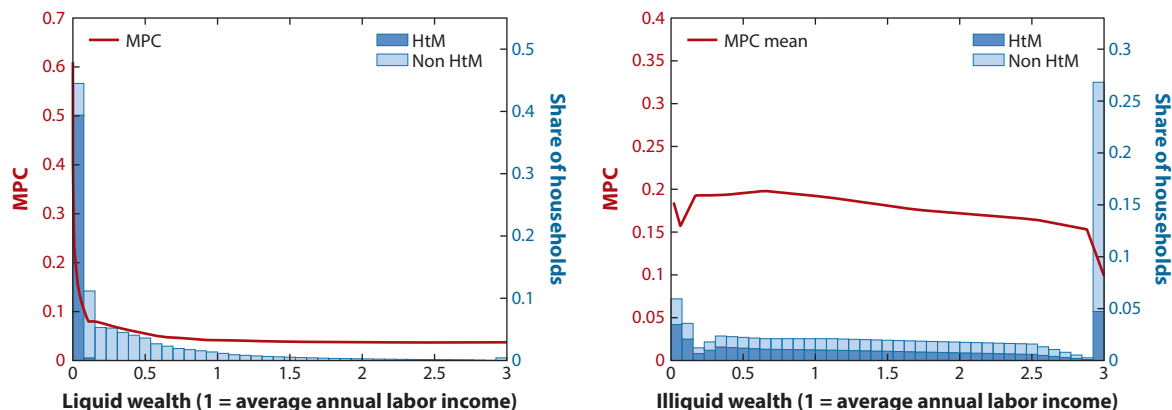


Figure 4

The MPC (left scale) and the household wealth distribution (right scale) for the baseline two-asset model as a function of liquid (*left*) and illiquid (*right*) wealth. Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume.

between the liquid and illiquid returns. We return to this feature of the calibration in the next section. The calibrated transaction cost is around \$1,500. On average, 9% of households rebalance their portfolios each quarter. Although the model generates less aggregate liquid wealth than in the data (0.23 versus 0.56), median liquid wealth matches almost exactly its empirical value, around 5% of average annual labor income. The empirical liquid wealth distribution is extremely right-skewed, and there is no force in the model that can deliver this feature of the data.¹³

4.1.3. Marginal propensities to consume in the two-asset model. The average quarterly MPC in the baseline two-asset model is 16.1% (Table 5, column 2), around five times larger than the corresponding average MPC in the continuous-time one-asset model (Table 5, column 1). A simple back-of-the-envelope calculation helps in rationalizing this magnitude. Multiplying the average MPC for poor HtM households (24%) by their share (13%), the average MPC for wealthy HtM households (30%) by their share (27%), and the average MPC for non-HtM households (7%) by their share (60%), and adding up these three components, we obtain an average MPC close to 16%. Figure 4 (*left*) illustrates the average MPCs as a function of liquid wealth holdings.

To shed some light on the sources of the higher average MPC in the two-asset model compared to the one-asset model, we report the results of the decomposition in Equation 4 in Table 5, panel C. We implement the decomposition by computing the average MPC in the two-asset model as a function of net worth w —that is, $m_0(w, y) = \int_{a+b=w} m_0(a, b, y) d\mu(a, b, y)$. The shapes of the consumption function and the wealth distribution play similar roles. Because of the presence of a high-return asset, the intertemporal saving motive is strong and the model asks for a lower discount factor to match total net worth, a force that increases the MPC at all wealth levels. In addition, the larger share of HtM consumers with low liquid wealth shifts the distribution toward the steep region of the MPC function.

Figure 4 (*left*) confirms that what accounts for the bulk of cross-sectional variation in MPCs is liquid wealth, in line with much of the recent empirical evidence discussed in the Introduction. The MPC function is instead quite flat as a function of illiquid wealth (Figure 4, *right*) and net

¹³This feature, however, is not crucial in determining the aggregate MPC because once households have enough liquid assets, their MPC converges to the certainty level (approximately equal to the discount rate).

worth, because there is a sizable share of wealthy HtM types even among households with high levels of illiquid wealth (up to around three times the annual income).

The statistics on the wealth distribution in **Table 5**, panel D, reveal that the two-asset model does not suffer from the missing-middle problem. Median net worth is 1.1, which is much closer to the empirical target of 1.56. The reason is that in the two-asset models households in the middle of the distribution hold the bulk of their wealth in illiquid assets, which have only a small effect on the MPC.

4.2. Alternative Calibrations of the Two-Asset Model

We now examine the robustness of our findings for the two-asset model with respect to some of the key parameters of the model. In these experiments, whenever we change parameters, we recalibrate only the discount rate ρ to match the same target for total wealth, unless otherwise noted.

4.2.1. Rates of return. Changing the rates of return on liquid and illiquid wealth has a significant impact on the MPCs, as illustrated in columns 3 to 6 of **Table 5**. Raising the liquid rate or reducing the illiquid rate narrows the return gap between the two assets, which leads to fewer HtM households, lowering the average MPC.¹⁴

4.2.2. Rebalancing frequency. **Supplemental Table E.7** illustrates that decreasing the frequency of rebalancing opportunities from an average of once per month to an average of once per quarter or once per year has only a small effect on the average MPC. However, these calibrations with less frequent rebalancing opportunities are more successful in matching the mean level of liquid wealth in the data (0.56). Households hold a larger buffer of liquid wealth both because opportunities to rebalance arrive less frequently and because households that are accumulating wealth must hold their savings in the liquid account for a longer duration, on average, before transferring them to their illiquid account. The model with rebalancing opportunities on average once per year also features a slightly larger average MPC, despite having fewer HtM households, because the less frequent adjustment opportunities mean that the wealthy HtM households that are waiting to withdraw resources have a higher MPC than in the baseline calibration (38% versus 30%).

4.2.3. Rebalancing cost. Varying the transaction cost from \$0 to \$3,000, while keeping the arrival rate of rebalancing opportunities at its baseline value of 3, has only very small effects on the average MPC, because there are offsetting effects on the distribution and the shape of the consumption function (see **Supplemental Table E.7**, columns 4–7).

4.3. Tackling the Large Return Differential

The baseline two-asset model can simultaneously match the evidence on MPCs and wealth targets only with a rate of return differential of around 8% per year. A reasonable question is whether a return differential of this magnitude is plausible, and if not, whether there are extensions to the model that could yield the same MPC and wealth as the statistics with a smaller return difference.

For the United States, from 1950 on, Jordà et al. (2019) estimate a gap of 5–6% between the real financial return on stocks and housing (which are the biggest components of illiquid wealth) versus the return on T-Bills. However, these calculations do not account for the potentially large

¹⁴The MPC decompositions in this section are performed relative to the baseline two-asset model in **Table 5**, column 2.

tax advantages from investing in housing due to deductibility of mortgage interest and favorable capital gains treatment, nor for the tax advantages that arise from investing in equities inside qualified retirement accounts. Moreover, in the United States many 401(k) plans include a generous employer matching contribution (often around 50%) that further raises the effective rate of return. Finally, owner-occupied housing yields a direct utility flow, which in previous work we estimated to be equivalent to a financial return of 3–4% per year (Kaplan & Violante 2014). Including any of these features in the model can lower the nominal return gap without changing the implied consequences for saving and spending patterns.

An alternative modification to reduce the rate of return differential is to add a commitment value of holding illiquid assets. This approach is similar in spirit to including a direct utility flow, because it amounts to modifying preferences in such a way that households desire to hold more illiquid assets as a fraction of total wealth for a given financial return gap. To assess the quantitative reach of this mechanism, in **Table 6** we introduce temptation and self-control, like in Section 3.2.1, into the two-asset model, along the lines proposed by Attanasio et al. (2020). Details of the model are in **Supplemental Appendix C.2**.

When we recalibrate the discount rate to keep mean total wealth constant at 4.1, but hold the return differential and adjustment cost fixed at their baseline values, adding even a small amount of temptation raises the aggregate MPC substantially. For example, with a temptation parameter of $\varphi = 0.01$, the average quarterly MPC rises from 16% to 24% (**Table 6**, column 2), and with a temptation parameter of $\varphi = 0.05$, the MPC rises to 47% (column 3). However, as in the one-asset versions of the model, these larger MPCs arise from having a much larger fraction of HtM households than in the data. The decomposition relative to the baseline two-asset model confirms that the majority of the higher MPC is due to the wealth distribution rather than the consumption function.

We can bring the two-asset model with temptation in line with the evidence on HtM households by also recalibrating the return gap and adjustment cost to target the fraction of wealthy HtM (41%) and poor HtM (14%) households. Columns 3 and 4 report the results of these calibrations with temptation parameters of $\varphi = 0.01$ and $\varphi = 0.015$, respectively. Both versions are successful in matching these targets, and both yield even higher average MPCs than in the baseline model, with much lower return differences. For example, with $\varphi = 0.015$ (column 5), the return gap is 4%, compared with 8% in the baseline, and the average MPC is 21%. In order to match the HtM targets with a lower return difference, the models with temptation require larger adjustment costs than in the baseline model.

5. OTHER CONCEPTS OF MARGINAL PROPENSITY TO CONSUME

So far, we have focused on the quarterly MPC out of an unexpected \$500 windfall. In this section, we extend the analysis to (a) MPCs out of unexpected income shocks of different sizes, both positive and negative; (b) MPCs at different horizons, including MPCs out of the news of future windfalls; and (c) MPCs out of unexpected illiquid injections (or illiquid capital gains) in the two-asset models.

5.1. Sign and Size Asymmetries

To explore sign and size asymmetries, we compute MPCs out of windfalls of $-\$500$ and $+\$5,000$.¹⁵ **Table 7** summarizes these results. Because of the concavity in the consumption function, the

¹⁵ Simulating MPCs out of large negative shocks has the problem that they lead to an empty budget set for many poor households, so one has to make ad hoc assumptions about what happens in that case.

Table 6 Two-asset model with temptation and self-control

	(1) Baseline two-asset	(2) $\varphi = 0.01$	(3) $\varphi = 0.05$	(4) $\varphi = 0.01$	(5) $\varphi = 0.015$
Rebalance arrival rate	3.0	3.0	3.0	3.0	3.0
Quarterly MPC (%)	16.1	23.8	46.7	18.1	21.1
Annual MPC (%)	41.2	49.7	63.9	42.5	44.3
Quarterly PHtM MPC (%)	24.3	24.8	51.9	23.9	23.6
Quarterly WHtM MPC (%)	29.8	47.8	75.9	40.3	50.7
Mean MPC at mean wealth (%)	7.0	10.4	37.0	6.5	13.3
Prob. HtM status at year t and year $t + 1$	0.9	0.9	0.9	0.9	0.9
Panel A: Calibrated variables					
Annualized discount factor	0.937	0.937	0.937	0.964	0.973
Annualized liquid return	−2.0	−2.0	−2.0	−2.0	−2.0
Annualized illiquid return	6.2	6.2	6.2	3.2	2.2
Rebalance cost (\$)	1,456.8	1,456.8	1,456.8	2,400.0	2,700.0
Panel B: Targeted statistics					
Mean total wealth	4.1	4.1	4.1	4.1	4.1
Share HtM	39.5	48.6	62.9	39.0	40.4
Share PHtM	13.2	18.0	24.0	13.0	13.3
Panel C: Decomposition					
Gap with baseline MPC	NA	7.7	30.5	2.0	5.0
Effect of MPC function	NA	0.9	5.0	−1.1	−1.5
Distributional effect	NA	2.0	15.1	0.1	0.2
Interaction	NA	4.7	10.4	3.0	6.3
Panel D: Wealth statistics					
Mean liquid wealth	0.2	0.2	0.1	0.2	0.2
Median total wealth	1.1	1.0	0.7	1.0	1.0
Median liquid wealth	0.1	0.0	0.0	0.1	0.0
$b \leq \$1,000$	33.0	44.3	60.2	36.0	38.2
$b \leq \$5,000$	54.7	59.7	70.2	55.3	56.4
$b \leq \$10,000$	66.1	69.8	78.8	66.2	67.1
$w \leq \$1,000$	10.0	14.2	23.4	10.3	10.5
$w \leq \$5,000$	20.8	22.4	26.0	24.0	24.5
$w \leq \$10,000$	23.5	24.6	27.5	29.1	30.0
$w \leq \$50,000$	45.9	47.2	51.0	46.4	47.2
$w \leq \$100,000$	57.7	58.8	61.3	58.0	58.4
Wealth, top 10% share	62.0	63.5	65.8	60.9	60.6
Wealth, top 1% share	19.2	20.2	21.6	16.2	15.2
Gini coefficient, total wealth	0.8	0.8	0.8	0.8	0.8

Column 1 is the baseline two-asset model. Columns 2–5 are two-asset models with Gul-Pesendorfer preferences solved with different temptation parameters. b denotes liquid wealth and $w = a + b$ denotes total wealth (or net worth). Abbreviations: HtM, hand-to-mouth; MPC, marginal propensity to consume; NA, not applicable; PHtM, poor hand-to-mouth; Prob., probability; WHtM, wealthy hand-to-mouth.

Table 7 Quarterly MPC for positive and negative windfalls of different sizes across various models

	(1) Baseline one-asset	(2) Mean(a) = 0.56	(3) β Het	(4) IES Het	(5) S-S	(6) Baseline two-asset
MPC (%) out of $-\$500$	5.0	14.9	19.5	22.7	17.2	26.6
MPC (%) out of $+\$500$	4.6	14.0	18.6	20.8	16.8	16.1
MPC (%) out of $+\$5,000$	3.6	11.3	15.4	14.5	13.2	12.3

Column 1 is the baseline one-asset model. Column 2 is the one-asset model calibrated to median net worth. Columns 3 and 4 are one-asset models with preference heterogeneity. Column 5 is the spender-saver model. Column 6 is the baseline two-asset model. Abbreviations: Het, heterogeneity; IES, intertemporal elasticity of substitution; MPC, marginal propensity to consume; S-S, spender-saver.

MPC out of small windfalls is bigger than the MPC out of larger windfalls. Concavity also implies that MPCs out of income losses are larger than MPCs out of income gains of the same absolute magnitude.

In the baseline one-asset model these effects are not very strong, because the asymmetries are significant only in regions in which the consumption function is very concave, where there is only a small share of households in the stationary distribution (Table 6, column 1). But in models with a larger share of households in regions of the state space where the consumption function is steep—such as the one-asset model calibrated to mean liquid wealth calibration, one-asset models with ex-ante heterogeneity, and two-asset models—there is stronger size and sign asymmetry (columns 2–7).

So far we have focused on MPCs out of a fixed dollar amount as opposed to a fixed fraction of each household's earnings. We adopted this measure as a baseline because there exists reliable quasi-experimental evidence that can be used to assess different models and because of the size asymmetries documented above. However, the average MPC out of an unexpected common percentage change in income is also interesting because it corresponds to the partial equilibrium impact of transitory aggregate income shocks with even incidence across all households. We computed the aggregate MPC out of a 1% increase in earnings for each household (mean earnings is \$67,000, so this corresponds to a mean windfall of \$670, which is close to \$500). For the one-asset model, the aggregate MPC out of this shock (computed as the total change in consumption divided by the total change in income) is 2.7%, which is close to the income-weighted MPC out of a \$500 gain, 2.7%. For the two-asset model, the aggregate MPC out of the 1% earnings shock is 13.1%. In both models, the aggregate MPC out of the 1% earnings shock is lower than the MPC out of a \$500 gain because higher-income households have lower MPCs.

5.2. Intertemporal Marginal Propensities to Consume

In general equilibrium versions of heterogeneous agent models, the entire time path of MPCs matters for the extent of shock propagation (see Auclert et al. 2018). We compute the time profile of MPCs from horizon $t = -4$ (one year before the shock) to $t = +4$ (one year after the shock). The quarterly MPC out of a \$500 gain at horizon $t = -\tau$ should be interpreted as the fraction of the windfall consumed by households upon receiving the news that a windfall will be received in τ quarters' time. The MPC at horizon $t = +\tau$ is the fraction consumed in τ quarters after receipt of an unexpected windfall. **Supplemental Appendix D** describes how to compute MPCs at different horizons.

Figure 5 reports the time profile of MPCs across several models. In the baseline one-asset model the profile is very flat because of the low share of HtM households. In one-asset models

Supplemental Material >

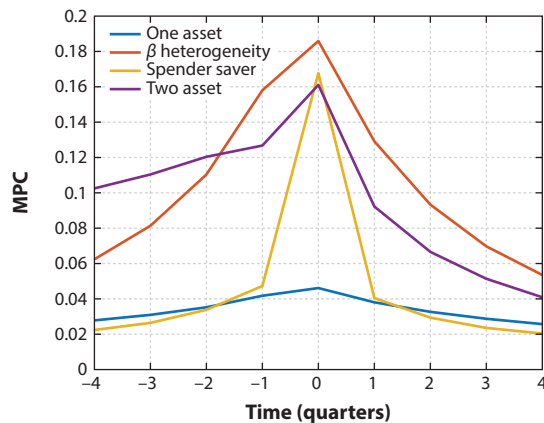


Figure 5

Intertemporal MPCs for four models. At $t = 0$ we observe the impact MPC; at $t > 0$, the lagged MPC out of a windfall received at $t = 0$; at $t < 0$, the MPC out of the news of a windfall that will be received at $t = 0$. Abbreviations: β , discount factor; MPC, marginal propensity to consume.

with discount factor heterogeneity, the profile is more tent-shaped: MPCs out of news are much smaller than actual MPCs, as are lagged MPCs. This is also true in versions of the model with IES heterogeneity (not displayed). These dynamics are driven by impatient or high-IES households. As pointed out by Auclert et al. (2018), the tent shape is particularly extreme in the spender-saver model. In this model, the time profile is an average of the MPCs for the savers, which closely mimics the shape of the profile in the baseline one-asset model, and the MPCs for the spenders, who have a strong spending response on impact but not at other leads or lags.

The time path for the two-asset model shows a high response at impact because of the large share of HtM households, as well as a sizable reaction to news about future windfalls. The reason for this latter effect is the low calibrated discount factor, which raises the MPCs of unconstrained households, which are the ones that respond to future windfalls. Comparing these anticipation effects in the model with corresponding empirical evidence requires caution, because in the model future windfalls are received with certainty, whereas in reality there is always some degree of uncertainty about whether the windfalls will ultimately be received.

In Section 2.2, we determined that in the one-asset model the annual MPC is much smaller than four times the quarterly MPC. **Figure 5** clearly illustrates that this conclusion holds across models, because the time path of the intertemporal MPCs is—sharply, in some economies—declining over time after $t = 0$.

5.3. Marginal Propensity to Consume out of Illiquid Wealth

In the two-asset model, we can also compute the MPC out of a small injection of illiquid wealth. Conceptually, the closest empirical counterpart of this MPC would be an estimated MPC out of changes in housing wealth or out of unrealized capital gains/losses on equity in retirement accounts or private businesses. To put our analysis in the context of recent empirical contributions, we note that Carroll et al. (2011) estimate an average quarterly MPC out of housing wealth of around 2%. Mian et al. (2013) estimate it at 1.5%, and they uncover a great deal of heterogeneity with respect to income and leverage. Di Maggio et al. (2020) analyze unrealized capital gains in stock market wealth and conclude that the literature, including their own research, finds values between 0 and 2.5% for the average quarterly MPC. In our model, the average quarterly MPC

out of an illiquid windfall of \$500 is 1.4%, and the MPC out of an illiquid windfall of \$5,000 is 2.4%. As in the data, the average illiquid MPC is roughly an order of magnitude smaller than the liquid MPC.

6. CONCLUSIONS

MPCs are the most important feature of household spending behavior for macroeconomics. They play an essential role in determining the effects of changes in aggregate demand, the strength of the fiscal multiplier, the transmission mechanism of monetary policy, and the portfolio share of wealth held in risky assets. Empirical evidence suggests that in response to an unexpected windfall, the MPC is large on average and displays a wide degree of heterogeneity across households. However, canonical representative agent models fail to account for this evidence.

One-asset heterogeneous agent models with idiosyncratic income risk and incomplete markets generate an average MPC that is an order of magnitude bigger than in representative agent models, but still significantly smaller than in the data. Allowing for ex-ante heterogeneity across households or for certain behavioral biases in preferences can raise the average MPC to be in line with the data. However, this comes at the cost of misrepresenting key features of the wealth distribution. Most importantly, these models vastly overstate the fraction of households with low wealth.

Two-asset heterogeneous agent models with a liquid and an illiquid asset can resolve this tension when they feature a large enough return gap between the two assets. Theory and data suggest a number of plausible justifications for such return differentials.

DISCLOSURE STATEMENT

The authors are not aware of any affiliations, memberships, funding, or financial holdings that might be perceived as affecting the objectivity of this review.

ACKNOWLEDGMENT

We thank Chase Abram and Brian Livingston for outstanding research assistance.

LITERATURE CITED

- Aguiar MA, Bils M, Boar C. 2020. *Who are the hand-to-mouth?* NBER Work. Pap. 26643
- Andersen AL, Jensen AS, Johannesen N, Kreiner CT, Leth-Petersen S, Sheridan A. 2020. *How do households respond to job loss? Lessons from multiple high-frequency data sets.* CEBI Work. Pap. 12/20, Cent. Econ. Behav. Inequal., Univ. Copenhagen, Copenhagen, Den.
- Attanasio O, Kovacs A, Moran P. 2020. *Temptation and commitment: understanding hand-to-mouth behavior.* NBER Work. Pap. 27944
- Auclert A. 2019. Monetary policy and the redistribution channel. *Am. Econ. Rev.* 109(6):2333–67
- Auclert A, Rognlie M, Straub L. 2018. *The intertemporal Keynesian cross.* NBER Work. Pap. 25020
- Aydin D. 2022. Consumption response to credit expansions: evidence from experimental assignment of 45,307 credit lines. *Am. Econ. Rev.* 112(1):1–40
- Benhabib J, Bisin A. 2018. Skewed wealth distributions: theory and empirics. *J. Econ. Lit.* 56(4):1261–91
- Bilbiie FO. 2020. The new Keynesian cross. *J. Monet. Econ.* 114:90–108
- Blundell R, Pistaferri L, Preston I. 2008. Consumption inequality and partial insurance. *Am. Econ. Rev.* 98(5):1887–921
- Bunn P, Le Roux J, Reinold K, Surico P. 2018. The consumption response to positive and negative income shocks. *J. Monet. Econ.* 96:1–15

- Campbell JY, Mankiw NG. 1989. Consumption, income, and interest rates: reinterpreting the time series evidence. *NBER Macroecon. Annu.* 4:185–216
- Carroll CD. 2001. A theory of the consumption function, with and without liquidity constraints. *J. Econ. Perspect.* 15(3):23–45
- Carroll CD, Holm MB, Kimball MS. 2021. Liquidity constraints and precautionary saving. *J. Econ. Theory* 195:105276
- Carroll CD, Kimball MS. 1996. On the concavity of the consumption function. *Econometrica* 64(4):981–92
- Carroll CD, Kimball MS. 2001. *Liquidity constraints and precautionary saving*. NBER Work. Pap. 8496
- Carroll CD, Otsuka M, Slacalek J. 2011. How large are housing and financial wealth effects? A new approach. *J. Money Credit Bank* 43(1):55–79
- Carroll CD, Slacalek J, Tokunaka K, White MN. 2017. The distribution of wealth and the marginal propensity to consume. *Quant. Econ.* 8(3):977–1020
- Commault J. 2022. Does consumption respond to transitory shocks? Reconciling natural experiments and semistructural methods. *Am. Econ. J. Macroecon.* 14(2):96–122
- De Nardi M, Fella G. 2017. Saving and wealth inequality. *Rev. Econ. Dyn.* 26:280–300
- Deaton A. 1991. Saving and liquidity constraints. *Econometrica* 59(5):1221–48
- Di Maggio M, Kermani A, Majlesi K. 2020. Stock market returns and consumption. *J. Finance* 75(6):3175–219
- Epstein LG, Zin SE. 1991. Substitution, risk aversion, and the temporal behavior of consumption and asset returns: an empirical analysis. *J. Political Econ.* 99(2):263–86
- Fagereng A, Holm MB, Natvik GJ. 2021. MPC heterogeneity and household balance sheets. *Am. Econ. J. Macroecon.* 13(4):1–54
- Fuster A, Kaplan G, Zafar B. 2021. What would you do with \$500? Spending responses to gains, losses, news, and loans. *Rev. Econ. Stud.* 88(4):1760–95
- Ganong P, Jones D, Noel P, Greig F, Farrell D, Wheat C. 2020. *Wealth, race, and consumption smoothing of typical income shocks*. NBER Work. Pap. 27552
- Gelman M. 2021. What drives heterogeneity in the marginal propensity to consume? Temporary shocks versus persistent characteristics. *J. Monet. Econ.* 117:521–42
- Gelman M, Kariv S, Shapiro MD, Silverman D, Tadelis S. 2020. How individuals respond to a liquidity shock: evidence from the 2013 government shutdown. *J. Public Econ.* 189:103917
- Golosov M, Graber M, Mogstad M, Novgorodsky D. 2021. *How Americans respond to idiosyncratic and exogenous changes in household wealth and unearned income*. NBER Work. Pap. 29000
- Gul F, Pesendorfer W. 2001. Temptation and self-control. *Econometrica* 69(6):1403–35
- Havranek T, Sokolova A. 2020. Do consumers really follow a rule of thumb? Three thousand estimates from 144 studies say “probably not.” *Rev. Econ. Dyn.* 35:97–122
- Helpman E. 1981. Optimal spending and money holdings in the presence of liquidity constraints. *Econometrica* 49(6):1559–70
- Jappelli T, Pistaferri L. 2010. The consumption response to income changes. *Annu. Rev. Econ.* 2:479–506
- Jappelli T, Pistaferri L. 2014. Fiscal policy and MPC heterogeneity. *Am. Econ. J. Macroecon.* 6(4):107–36
- Johnson DS, Parker JA, Souleles NS. 2006. Household expenditure and the income tax rebates of 2001. *Am. Econ. Rev.* 96(5):1589–610
- Jordà Ò, Knoll K, Kuvshinov D, Schularick M, Taylor AM. 2019. The rate of return on everything, 1870–2015. *Q. J. Econ.* 134(3):1225–98
- Kaplan G, Moll B, Violante GL. 2018. Monetary policy according to hank. *Am. Econ. Rev.* 108(3):697–743
- Kaplan G, Violante GL. 2014. A model of the consumption response to fiscal stimulus payments. *Econometrica* 82(4):1199–239
- Kaplan G, Violante GL, Weidner J. 2014. The wealthy hand-to-mouth. *Brook. Pap. Econ. Activ.* Spring:77–138
- Kekre R, Lenel M. 2021. *Monetary policy, redistribution, and risk premia*. NBER Work. Pap. 28869
- Kimball MS. 1990. *Precautionary saving and the marginal propensity to consume*. NBER Work. Pap. 8233
- Krueger D, Mitman K, Perri F. 2016. Macroeconomics and household heterogeneity. In *Handbook of Macroeconomics*, Vol. 2A, ed. JB Taylor, H Uhlig, pp. 843–921. Amsterdam: Elsevier
- Krusell P, Kuruşçu B, Smith AA Jr. 2002. Time orientation and asset prices. *J. Monet. Econ.* 49(1):107–35
- Krusell P, Kuruşçu B, Smith AA Jr. 2010. Temptation and taxation. *Econometrica* 78(6):2063–84

- Kueng L. 2018. Excess sensitivity of high-income consumers. *Q. J. Econ.* 133(4):1693–751
- Laibson D. 1997. Golden eggs and hyperbolic discounting. *Q. J. Econ.* 112(2):443–78
- Laibson D, Maxted P, Moll B. 2021. *Present bias amplifies the household balance-sheet channels of macroeconomic policy*. NBER Work. Pap. 29094
- Lewis DJ, Melcangi D, Pilossoph L. 2019. *Latent heterogeneity in the marginal propensity to consume*. Staff Rep. 902, Fed. Reserve Bank New York, New York
- Mian A, Rao K, Sufi A. 2013. Household balance sheets, consumption, and the economic slump. *Q. J. Econ.* 128(4):1687–726
- Misra K, Surico P. 2014. Consumption, income changes, and heterogeneity: evidence from two fiscal stimulus programs. *Am. Econ. J. Macroecon.* 6(4):84–106
- Nakajima M. 2017. Assessing bankruptcy reform in a model with temptation and equilibrium default. *J. Public Econ.* 145:42–64
- Parker JA, Souleles NS. 2019. Reported effects versus revealed-preference estimates: evidence from the propensity to spend tax rebates. *Am. Econ. Rev. Insights* 1(3):273–90
- Parker JA, Souleles NS, Johnson DS, McClelland R. 2013. Consumer spending and the economic stimulus payments of 2008. *Am. Econ. Rev.* 103(6):2530–53
- Patterson C. 2019. *The matching multiplier and the amplification of recessions*. Work. Pap., Northwest. Univ., Evanston, IL
- Pavoni N, Yazici H. 2017. Optimal life-cycle capital taxation under self-control problems. *Econ. J.* 127(602):1188–216
- Shapiro MD, Slemrod J. 2003. Consumer response to tax rebates. *Am. Econ. Rev.* 93(1):381–96
- Shapiro MD, Slemrod J. 2009. Did the 2008 tax rebates stimulate spending? *Am. Econ. Rev.* 99(2):374–79
- Werning I. 2015. *Incomplete markets and aggregate demand*. NBER Work. Pap. 21448
- Wolf C. 2020. *The missing intercept: a demand equivalence approach*. Tech. Rep., Mass. Inst. Technol., Cambridge
- Zeldes SP. 1989. Consumption and liquidity constraints: an empirical investigation. *J. Political Econ.* 97(2):305–



Contents

The Great Divide: Education, Despair, and Death <i>Anne Case and Angus Deaton</i>	1
The Impact of Health Information and Communication Technology on Clinical Quality, Productivity, and Workers <i>Ari Bronsoler, Joseph Doyle, and John Van Reenen</i>	23
Household Financial Transaction Data <i>Scott R. Baker and Lorenz Kueng</i>	47
Media and Social Capital <i>Filipe Campante, Ruben Durante, and Andrea Tesei</i>	69
The Elusive Explanation for the Declining Labor Share <i>Gene M. Grossman and Ezra Oberfield</i>	93
The Past and Future of Economic Growth: A Semi-Endogenous Perspective <i>Charles I. Jones</i>	125
Risks and Global Supply Chains: What We Know and What We Need to Know <i>Richard Baldwin and Rebecca Freeman</i>	153
Managing Retirement Incomes <i>James Banks and Rowena Crawford</i>	181
The Economic Impacts of the US–China Trade War <i>Pablo D. Fajgelbaum and Amit K. Khandelwal</i>	205
How Economic Development Influences the Environment <i>Seema Jayachandran</i>	229
The Economics of the COVID-19 Pandemic in Poor Countries <i>Edward Miguel and Ahmed Mushfiq Mobarak</i>	253
The Affordable Care Act After a Decade: Industrial Organization of the Insurance Exchanges <i>Benjamin Handel and Jonathan Kolstad</i>	287
Helicopter Money: What Is It and What Does It Do? <i>Ricardo Reis and Silvana Tenreyro</i>	313

Relational Contracts and Development <i>Rocco Macchiavello</i>	337
Trade Policy Uncertainty <i>Kyle Handley and Nuno Limão</i>	363
Bureaucracy and Development <i>Timothy Besley, Robin Burgess, Adnan Khan, and Guo Xu</i>	397
Misperceptions About Others <i>Leonardo Bursztyn and David Y. Yang</i>	425
The Affordable Care Act After a Decade: Its Impact on the Labor Market and the Macro Economy <i>Hanming Fang and Dirk Krueger</i>	453
Expecting Brexit <i>Swati Dhingra and Thomas Sampson</i>	495
Salience <i>Pedro Bordalo, Nicola Gennaioli, and Andrei Shleifer</i>	521
Enough Potential Repudiation: Economic and Legal Aspects of Sovereign Debt in the Pandemic Era <i>Anna Gelpern and Ugo Panizza</i>	545
The Great Gatsby Curve <i>Steven N. Durlauf, Andros Kourtellos, and Chih Ming Tan</i>	571
Inequality and the COVID-19 Crisis in the United Kingdom <i>Richard Blundell, Monica Costa Dias, Jonathan Cribb, Robert Joyce, Tom Waters, Thomas Wernham, and Xiaowei Xu</i>	607
The Aftermath of Debt Surges <i>M. Ayhan Kose, Franziska L. Ohnsorge, Carmen M. Reinhart, and Kenneth S. Rogoff</i>	637
Networks and Economic Fragility <i>Matthew Elliott and Benjamin Golub</i>	665
Central Bank Digital Currencies: Motives, Economic Implications, and the Research Frontier <i>Raphael Auer, Jon Frost, Leonardo Gambacorta, Cyril Monnet, Tara Rice, and Hyun Song Shin</i>	697
The Use of Scanner Data for Economics Research <i>Pierre Dubois, Rachel Griffith, and Martin O'Connell</i>	723
The Marginal Propensity to Consume in Heterogeneous Agent Models <i>Greg Kaplan and Giovanni L. Violante</i>	747

Experimental Economics: Past and Future <i>Guillaume R. Fréchet, Kim Sarnoff, and Leeat Yariv</i>	777
Spatial Sorting and Inequality <i>Rebecca Diamond and Cecile Gaubert</i>	795
Regression Discontinuity Designs <i>Matias D. Cattaneo and Rocío Titiunik</i>	821
Early Childhood Development, Human Capital, and Poverty <i>Orazio Attanasio, Sarah Cattan, and Costas Meghir</i>	853
The Econometric Model for Causal Policy Analysis <i>James J. Heckman and Rodrigo Pinto</i>	893

Indexes

Cumulative Index of Contributing Authors, Volumes 10–14	925
---	-----

Errata

An online log of corrections to *Annual Review of Economics* articles may be found at
<http://www.annualreviews.org/errata/economics>